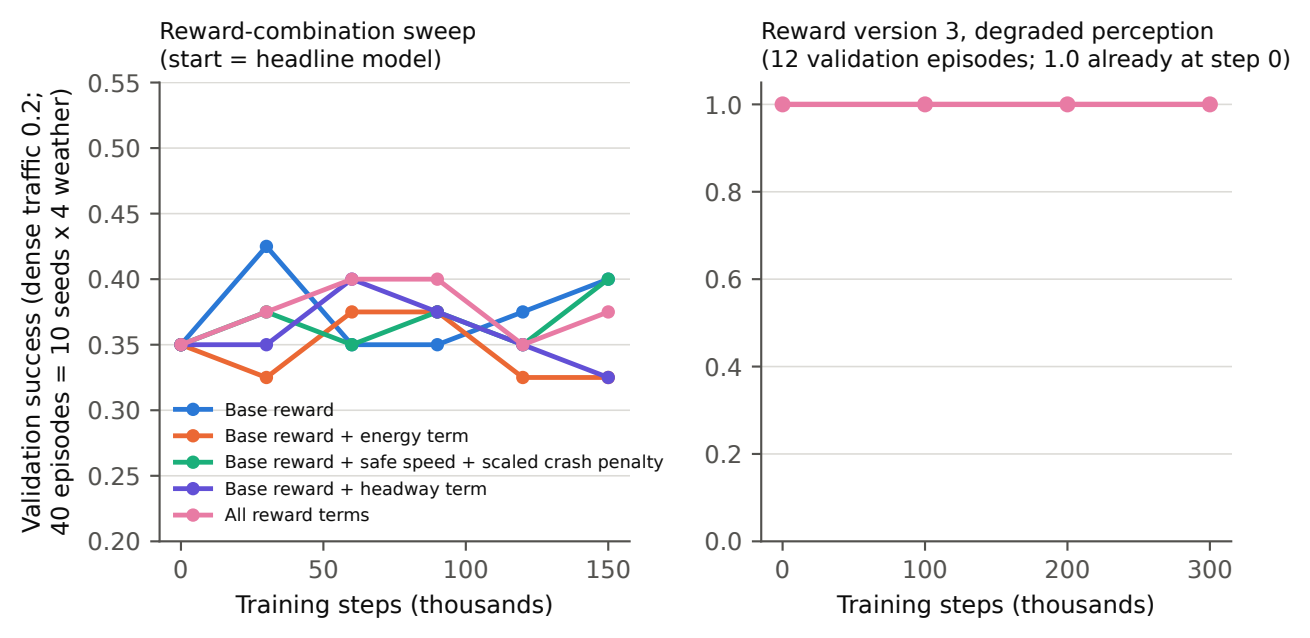
Reward-combination sweep
(start = headline model)
Validation success (dense traffic 0.2;
40 episodes = 10 seeds x 4 weather)
0.55
0.50
0.45
0.40
0.35
0.30
0.25
0.20
0
50
100
150
Training steps (thousands)
Base reward
Base reward + energy term
Base reward + safe speed + scaled crash penalty
Base reward + headway term
All reward terms
Reward version 3, degraded perception
(12 validation episodes; 1.0 already at step 0)
1.0
0.8
0.6
0.4
0.2
0.0
0
100
200
300
Training steps (thousands)